

Case Study 2:

Case Study Description:

Let us take up the CUSTOMER and TRANSACTIONS table we have created in the Lets Do Together section. Let us solve the following use cases using these tables :-

As part of the case study, Creating 2 tables namely customer and transaction using hive and loading the data from a txt file. After doing datanode format all the data is cleared out.

```
hive> CREATE TABLE CUSTOMER(  
  > custid INT,  
  > fname STRING,  
  > lname STRING,  
  > age INT,  
  > profession STRING  
  > )  
  > row format delimited fields terminated by ',';  
OK  
Time taken: 0.698 seconds  
hive> LOAD DATA LOCAL INPATH '/home/acadgild/custs.txt' into table CUSTOMER;  
Loading data to table default.customer  
OK  
Time taken: 1.357 seconds  
hive> CREATE TABLE TXNRECORDS (txnno INT, txndate STRING, custno INT, amount DOUBLE, category STRING, product STRING, city STRING  
, state STRING, spendby STRING)  
  > row format delimited fields terminated by ',';  
OK  
Time taken: 0.066 seconds  
hive> LOAD DATA LOCAL INPATH '/home/acadgild/txns.txt' into table TXNRECORDS;  
Loading data to table default.txnrecords  
OK  
Time taken: 0.341 seconds
```

```
hive> select * from customer;  
OK  
4000001 Kristina Chung 55 Pilot  
4000002 Paige Chen 74 Teacher  
4000003 Sherri Melton 34 Firefighter  
4000004 Gretchen Hill 66 Computer hardware engineer  
4000005 Karen Puckett 74 Lawyer  
4000006 Patrick Song 42 Veterinarian  
4000007 Elsie Hamilton 43 Pilot  
4000008 Hazel Bender 63 Carpenter  
4000009 Malcolm Wagner 39 Artist  
4000010 Dolores McLaughlin 60 Writer  
Time taken: 1.535 seconds, Fetched: 10 row(s)  
hive> select * from txnrecords;  
OK  
0 06-26-2011 4000001 40.33 Exercise & Fitness Cardio Machine Accessories Clarksville Tennessee c  
redit  
1 05-26-2011 4000002 198.44 Exercise & Fitness Weightlifting Gloves Long Beach California credit  
2 06-01-2011 4000002 5.58 Exercise & Fitness Weightlifting Machine Accessories Anaheim California c  
redit  
3 06-05-2011 4000003 120.38 Exercise & Fitness Cardio Machine Accessories Clarksville Tennessee c  
redit  
4 06-05-2011 4000003 120.38 Exercise & Fitness Cardio Machine Accessories Clarksville Tennessee c  
redit  
5 06-05-2011 4000003 120.38 Exercise & Fitness Cardio Machine Accessories Clarksville Tennessee c  
redit  
6 06-05-2011 4000003 120.38 Exercise & Fitness Cardio Machine Accessories Clarksville Tennessee c  
redit  
7 06-05-2011 4000003 120.38 Exercise & Fitness Cardio Machine Accessories Clarksville Tennessee c  
redit  
8 06-05-2011 4000003 120.38 Exercise & Fitness Cardio Machine Accessories Clarksville Tennessee c  
redit  
9 06-05-2011 4000003 120.38 Exercise & Fitness Cardio Machine Accessories Clarksville Tennessee c  
redit
```

1. Find out the number of transaction done by each customer (These should be take up in module 8 itself)

Query:

```
select a.custid, a.fname, count(b.custno) from CUSTOMER a join  
TXNRECORDS b on a.custid = b.custno group by a.custid, a.fname;
```

Output:

```
4000001 Kristina      8  
4000002 Paige       6  
4000003 Sherri      3  
4000004 Gretchen    5  
4000005 Karen       5  
4000006 Patrick     5  
4000007 Elsie       6  
4000008 Hazel      10  
4000009 Malcolm     6  
4000010 Dolores     6  
Time taken: 24.788 seconds, Fetched: 10 row(s)
```

2. Create a new table called TRANSACTIONS_COUNT. This table should have 3 fields - custid, fname and count. (Again to be done in module 8)

```
CREATE TABLE TRANSACTIONS_COUNT(custid INT, fname  
STRING,count INT) row format delimited fields terminated by ',';
```

Output:

```
hive> CREATE TABLE TRANSACTIONS_COUNT(  
  > custid INT,  
  > fname STRING,count INT  
  > )  
  > row format delimited fields terminated by ',';  
OK  
Time taken: 0.302 seconds  
hive> select * from transactions_count;  
OK  
Time taken: 0.097 seconds  
hive> █
```

3. Now write a hive query in such a way that the query populates the data obtained in Step 1 above and populate the table in step 2 above. (This has to be done in module 9).

Query:

Insert into transaction_count select a.custid, a.fname, count(b.custno) from CUSTOMER a join TXNRECORDS b on a.custid = b.custno group by a.custid, a.fname;

Output:

```
hive> Insert into transaction_count select a.custid, a.fname, count(b.custno) from CUSTOMER a join TXNRECORDS b on a.custid = b.custno group by a.custid, a.fname;
FAILED: SemanticException org.apache.hadoop.hive.ql.metadata.InvalidTableException: Table not found transaction_count
hive> Insert into transactions_count select a.custid, a.fname, count(b.custno) from CUSTOMER a join TXNRECORDS b on a.custid = b.custno group by a.custid, a.fname;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180518115747_a0aecf56-b7a2-46d9-be60-c9e4023479e1
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-05-18 11:57:51 Starting to launch local task to process map join; maximum memory = 477626368
2018-05-18 11:57:52 Dump the side-table for tag: 0 with group count: 10 into file: file:/tmp/acadgild/867b4b5b-e7fc-44fb-8904-5555f7fa19f7/hive_2018-05-18_11-57-47_087_1684632309728826396-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile10--.hashtable
2018-05-18 11:57:52 Uploaded 1 File to: file:/tmp/acadgild/867b4b5b-e7fc-44fb-8904-5555f7fa19f7/hive_2018-05-18_11-57-47_087_1684632309728826396-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile10--.hashtable (556 bytes)
2018-05-18 11:57:52 End of local task; Time Taken: 1.167 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1526623851877_0002, Tracking URL = http://localhost:8088/proxy/application_1526623851877_0002/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1526623851877_0002
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
```

```
hive> select * from transactions_count;
OK
4000001 Kristina      8
4000002 Paige 6
4000003 Sherri 3
4000004 Gretchen    5
4000005 Karen 5
4000006 Patrick 5
4000007 Elsie 6
4000008 Hazel 10
4000009 Malcolm 6
4000010 Dolores 6
Time taken: 0.087 seconds, Fetched: 10 row(s)
```

4. Now let's make the TRANSACTIONS_COUNT table Hbase compliant. In the sense, use Ser Des And Storate handler features of hive to change the TRANSACTIONS_COUNT table to be able to create a TRANSACTIONS table in Hbase. (This has to be done in module 10)

Query:

```
-----  
CREATE TABLE TRANSACTIONS(custId int, fname string, count int)  
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'  
WITH SERDEPROPERTIES ("hbase.columns.mapping" =  
":key,custom_details:fname, custom_details:count")  
TBLPROPERTIES ("hbase.table.name" = "TRANSACTIONS");
```

Output:

```
-----  
hive> CREATE TABLE TRANSACTIONS(custId int, fname string, count int)  
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'  
> WITH SERDEPROPERTIES ("hbase.columns.mapping" = ":key,custom_details:fname, custom_details:count")  
> TBLPROPERTIES ("hbase.table.name" = "TRANSACTIONS");  
OK  
Time taken: 2.596 seconds
```

```
hbase(main):002:0> list  
TABLE  
TRANSACTIONS  
bulkdata  
clicks  
3 row(s) in 0.0550 seconds  
  
=> ["TRANSACTIONS", "bulkdata", "clicks"]  
hbase(main):003:0> scan 'TRANSACTIONS'  
ROW COLUMN+CELL  
0 row(s) in 0.0210 seconds
```

```
hbase(main):004:0> describe 'TRANSACTIONS'  
Table TRANSACTIONS is ENABLED  
TRANSACTIONS  
COLUMN FAMILIES DESCRIPTION  
{NAME => 'custom_details', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK  
ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', RE  
PLICATION_SCOPE => '0'}  
1 row(s) in 0.1480 seconds
```

5. Now insert the data in TRANSACTIONS_COUNT table using the query in step 3 again, this should populate the Hbase

TRANSACTIONS table automatically (This has to be done in module 10)

Query:

INSERT OVERWRITE table Transactions select a.custid, a.fname, count(b.custno) from CUSTOMER a join TXNRECORDS b on a.custid = b.custno group by a.custid, a.fname;

Output:

```
hbase(main):005:0> scan 'TRANSACTIONS'
ROW                                COLUMN+CELL
4000001                            column=custom_details:count, timestamp=1526628285357, value=8
4000001                            column=custom_details:fname, timestamp=1526628285357, value=Kristina
4000002                            column=custom_details:count, timestamp=1526628285357, value=6
4000002                            column=custom_details:fname, timestamp=1526628285357, value=Paige
4000003                            column=custom_details:count, timestamp=1526628285357, value=3
4000003                            column=custom_details:fname, timestamp=1526628285357, value=Sherri
4000004                            column=custom_details:count, timestamp=1526628285357, value=5
4000004                            column=custom_details:fname, timestamp=1526628285357, value=Gretchen
4000005                            column=custom_details:count, timestamp=1526628285357, value=5
4000005                            column=custom_details:fname, timestamp=1526628285357, value=Karen
4000006                            column=custom_details:count, timestamp=1526628285357, value=5
4000006                            column=custom_details:fname, timestamp=1526628285357, value=Patrick
4000007                            column=custom_details:count, timestamp=1526628285357, value=6
4000007                            column=custom_details:fname, timestamp=1526628285357, value=Elsie
4000008                            column=custom_details:count, timestamp=1526628285357, value=10
4000008                            column=custom_details:fname, timestamp=1526628285357, value=Hazel
4000009                            column=custom_details:count, timestamp=1526628285357, value=6
4000009                            column=custom_details:fname, timestamp=1526628285357, value=Malcolm
4000010                            column=custom_details:count, timestamp=1526628285357, value=6
4000010                            column=custom_details:fname, timestamp=1526628285357, value=Dolores
10 row(s) in 0.0530 seconds
```

```
hbase(main):009:0> get 'TRANSACTIONS', '4000001'
COLUMN                                CELL
custom_details:count                  timestamp=1526628285357, value=8
custom_details:fname                  timestamp=1526628285357, value=Kristina
2 row(s) in 0.0230 seconds
```

6. Now from the Hbase level, write the Hbase java API code to access and scan the TRANSACTIONS table data from java level (This should be done in module 11)

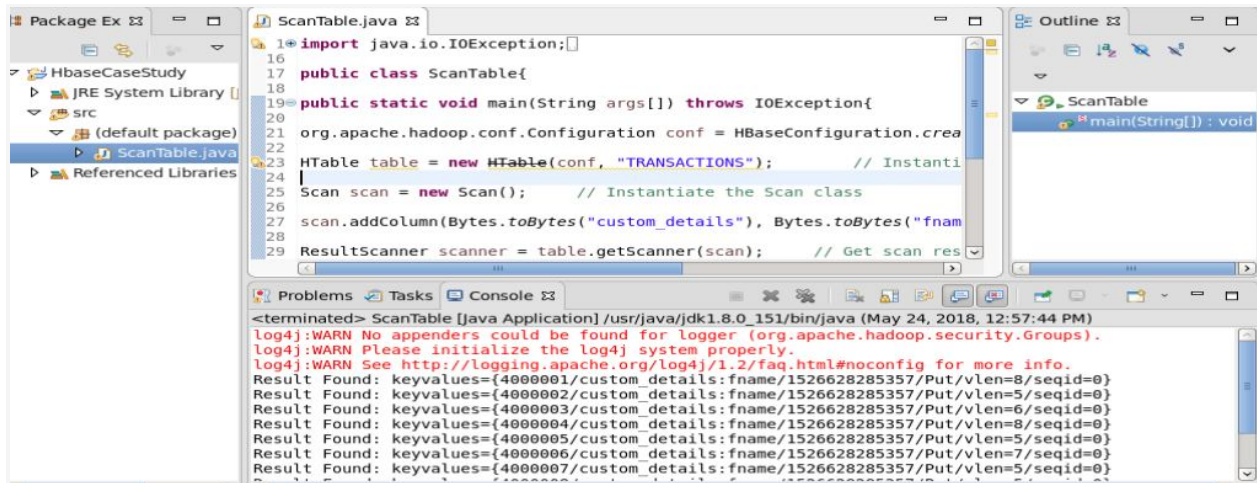
Steps:

Write a java API code and run as Java application in eclipse.

Executed in VM Machine.

The source code is uploaded as separate file.

Output:



The screenshot displays the Eclipse IDE interface. The Package Explorer on the left shows the project structure with 'ScanTable.java' selected. The main editor window shows the source code of 'ScanTable.java', which includes imports for 'java.io.IOException', 'org.apache.hadoop.conf.Configuration', 'org.apache.hadoop.hbase.HBaseConfiguration', 'org.apache.hadoop.hbase.HTable', 'org.apache.hadoop.hbase.Scan', and 'org.apache.hadoop.hbase.ResultScanner'. The code defines a 'main' method that creates an HTable, a Scan object, adds a column, and uses a ResultScanner to iterate through results. The Outline view on the right shows the 'main' method. The Console view at the bottom shows the execution output, including log4j warnings and the results of the scan.

```
1 import java.io.IOException;
16
17 public class ScanTable{
18
19 public static void main(String args[]) throws IOException{
20
21 org.apache.hadoop.conf.Configuration conf = HBaseConfiguration.crea
22
23 HTable table = new HTable(conf, "TRANSACTIONS"); // Instanti
24
25 Scan scan = new Scan(); // Instantiate the Scan class
26
27 scan.addColumn(Bytes.toBytes("custom_details"), Bytes.toBytes("fnam
28
29 ResultScanner scanner = table.getScanner(scan); // Get scan res
```

```
<terminated> ScanTable [Java Application] /usr/java/jdk1.8.0_151/bin/java (May 24, 2018, 12:57:44 PM)
log4j:WARN No appenders could be found for logger (org.apache.hadoop.security.Groups).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Result Found: keyvalues={4000001/custom_details:fname/1526628285357/Put/vlen=8/seqid=0}
Result Found: keyvalues={4000002/custom_details:fname/1526628285357/Put/vlen=5/seqid=0}
Result Found: keyvalues={4000003/custom_details:fname/1526628285357/Put/vlen=6/seqid=0}
Result Found: keyvalues={4000004/custom_details:fname/1526628285357/Put/vlen=8/seqid=0}
Result Found: keyvalues={4000005/custom_details:fname/1526628285357/Put/vlen=5/seqid=0}
Result Found: keyvalues={4000006/custom_details:fname/1526628285357/Put/vlen=7/seqid=0}
Result Found: keyvalues={4000007/custom_details:fname/1526628285357/Put/vlen=5/seqid=0}
```